

Generative Models

2024. 8

Yongjin Jeong, KwangWoon University

[참고] 본 자료에는 인터넷에서 다운받아 사용한 그림이나 수식들이 일부 있으니 다른 용도로 사용하거나 외부로 유출을 금해 주시기 바랍니다.

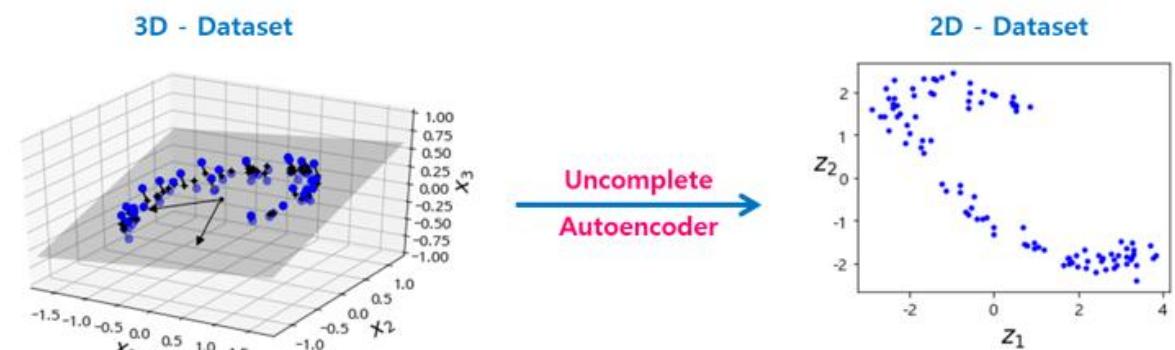
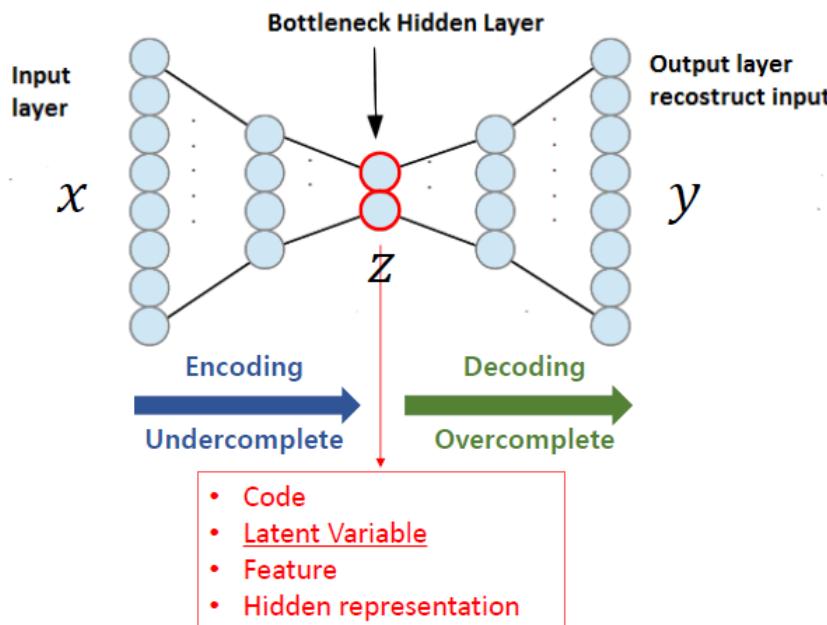
Contents

- Autoencoders
- Variational Autoencoder (VAE)
- GANs

Autoencoders

- **Autoencoder**

- To learn a representation (encoding) for training data set (**dimensionality reduction**) by training the network to ignore signal "noise", in other words, **feature extraction**.
- Focus on the **encoder**.



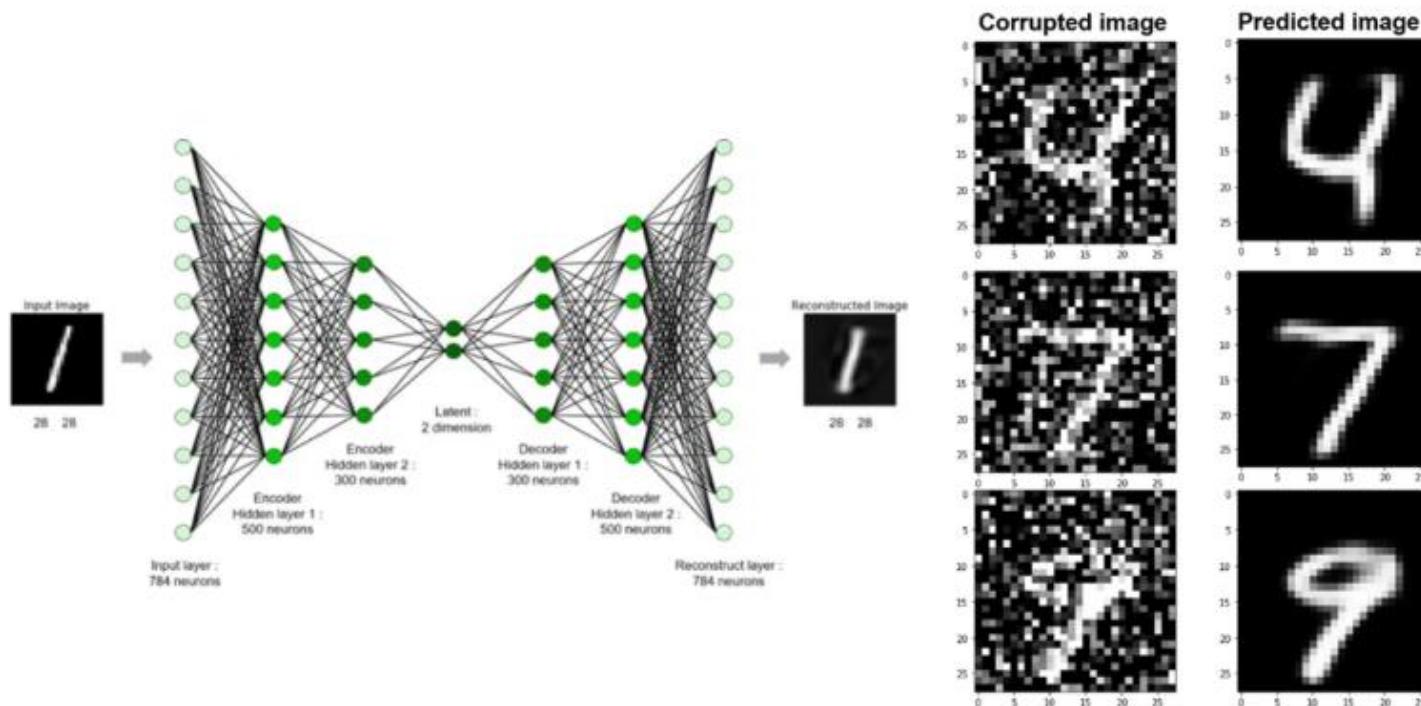
Autoencoders

- **Applications of Autoencoders**

- Dimensionality reduction
- Anomaly detection
- Image denoising
- Image compression
- Image generation

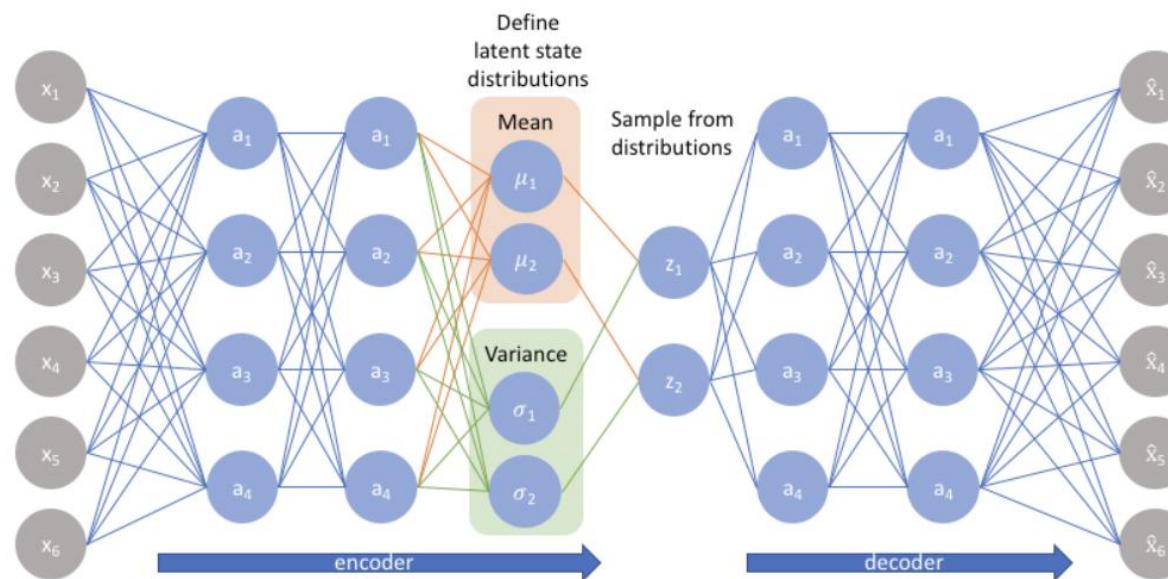
Denoising Autoencoder

- Denoising Autoencoder
 - Add noise to the input and learn to reconstruct the original input without noise
 - Noise can be generated in random input units using Gaussian Noise or Dropout.



Variational Autoencoder

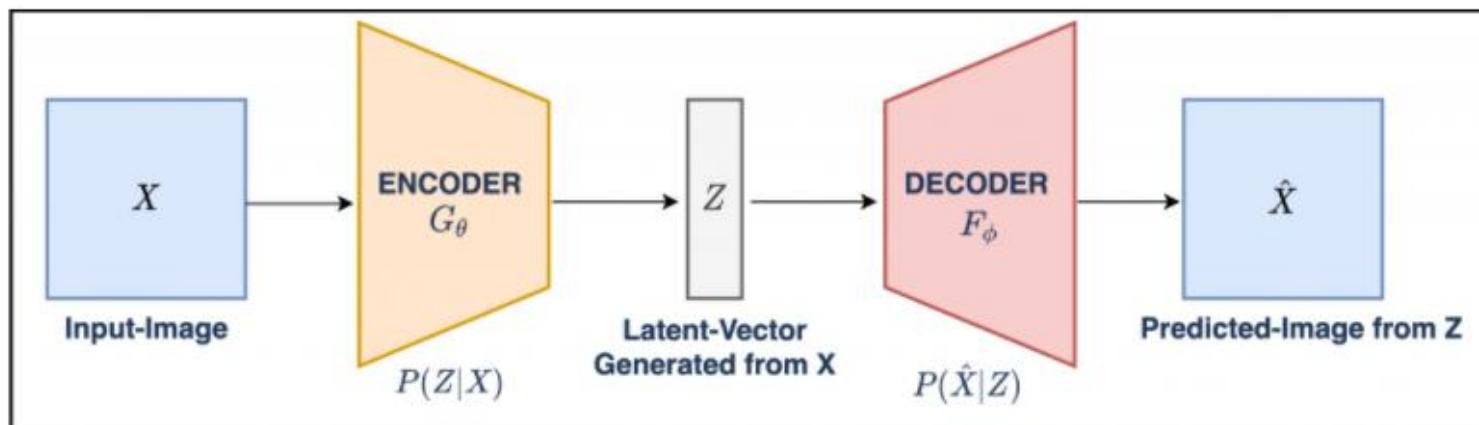
- **Variational Autoencoder (VAE)**
 - Provides a probabilistic manner for describing an observation in latent space
 - Encoder is formulated to describe a **probability distribution** for each latent attribute
 - Encoder is referred to as **recognition model**, and decoder is referred to as the **generative model**



Autoencoders

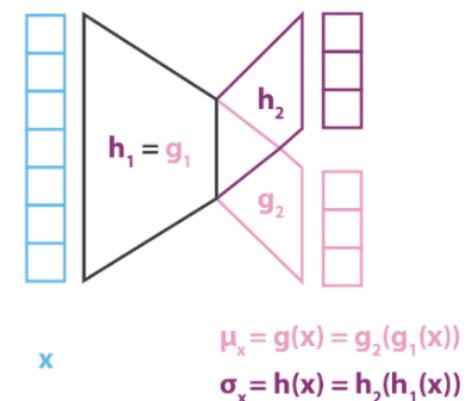
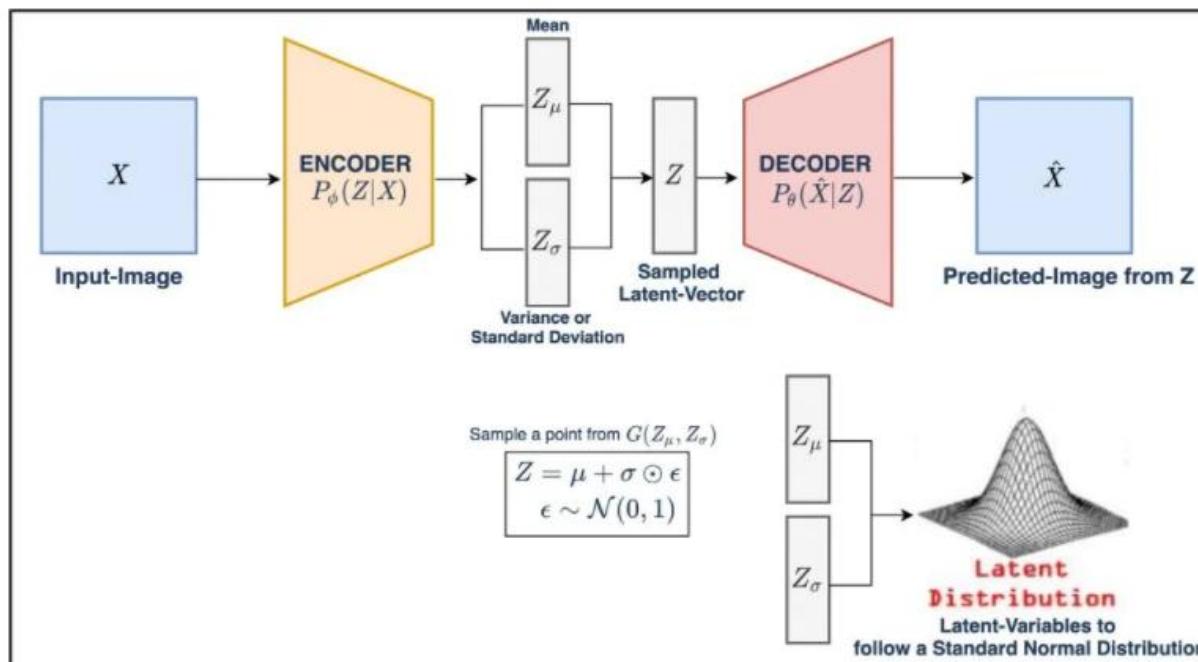
- **Auto-Encoder**

- **Reconstruction loss (MSE)** = $\frac{1}{N} \sum_{i=1}^N (X_i - \hat{X}_i)^2$ (N is number of images in a batch)
- It is not good at generating new images because of its latent space structure:
 - The latent space was **not continuous** and did not allow easy interpolation.
 - Encoded vectors are grouped in clusters corresponding to different data classes, and there are huge gaps between the clusters.
 - While generating a new sample, the decoder often produces a gibberish output if the chosen point in the latent space did not contain any data



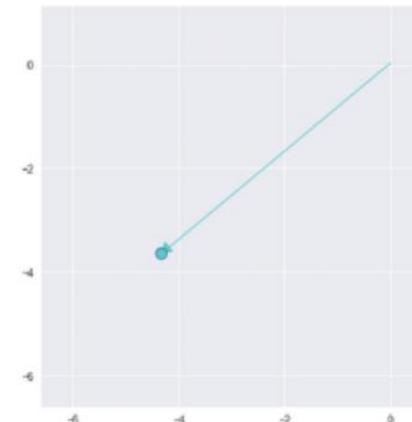
Variational Autoencoder (VAE)

- A **probabilistic** and **generative** model (focus on the **decoder**)
- Instead of mapping the image on a point in space, the encoder of VAE maps the image onto a **normal distribution**.
- Sample a **latent vector Z** from two **latent variables** Z_μ and Z_σ (also called a **sampling-layer**)
- We want the latent vector Z to follow a standard normal distribution. (Z_μ and Z_σ should be trained such that they are close to 0 and 1 respectively)

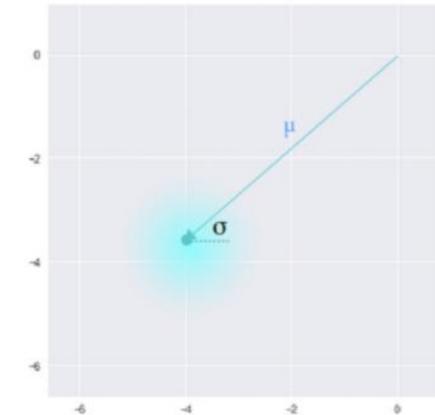


Variational Autoencoder (VAE)

- AE and VAE: Discrete latent vector and a standard normal distribution
 - AE: 연속성 보장이 안됨 – 생성 모델로 부적합
 - 단순히 입력을 압축한 벡터라서, 잠재 공간이 불연속적일 수 있다.
 - 특정 데이터 샘플 주변에만 의미 있는 벡터가 존재하고, 그 사이 공간은 의미 없는 값일 수 있다.
 - VAE: 연속성 보장
 - 잠재 공간을 정규분포에 맞추도록 KL 발산 손실을 추가.
 - 덕분에 잠재 공간 전체가 매끄럽게 연결되어 있어, 임의의 점을 샘플링해도 의미 있는 데이터가 생성.
- Why normal distribution?
 - 수학적 단순성
 - 재매개변수화 트릭(Reparameterization Trick)
 - 잠재 공간의 연속성
 - 샘플링 용이성



Standard Autoencoder
(direct encoding coordinates)



Variational Autoencoder
(μ and σ initialize a probability distribution)

Variational Autoencoder (VAE)

- **Objective Function of VAE**

- Unit Gaussian Distribution $\mathcal{N}(0, 1)$, and
- Minimize the reconstruction error $\frac{1}{N} \sum_{i=1}^N (X_i - \hat{X}_i)^2$

- **VAE's total loss**

- Loss = Re-construction_loss + KL_divergence_loss

$$L_{MSE}(\theta, \phi) = \frac{1}{N} \sum_{i=1}^N (x_i - f_\theta(g_\phi(x_i)))^2 \quad \text{or CE (Cross Entropy)}$$

$$L_{KL}[G(Z_\mu, Z_\sigma) \| \mathcal{N}(0, 1)] = -0.5 * \sum_{i=1}^N 1 + \log(Z_{\sigma_i}^2) - Z_{\mu_i}^2 - Z_{\sigma_i}^2$$

- MSE (continuous values): 입력 데이터가 연속적인 값을 가지는 경우 (예: 자연 이미지, 음성 신호)
- CE (discrete values): 입력 데이터가 0~1 범위의 픽셀 값을 가지는 경우 (예: MNIST 손글씨 이미지)
- 주로 Binary CE 를 사용함. (이 경우 y_i 값은 0,1 이 아닌 확률 기대값으로 해석) (예: $y_i=0.7$: 해당 픽셀이 "1" 일 확률이 70%)

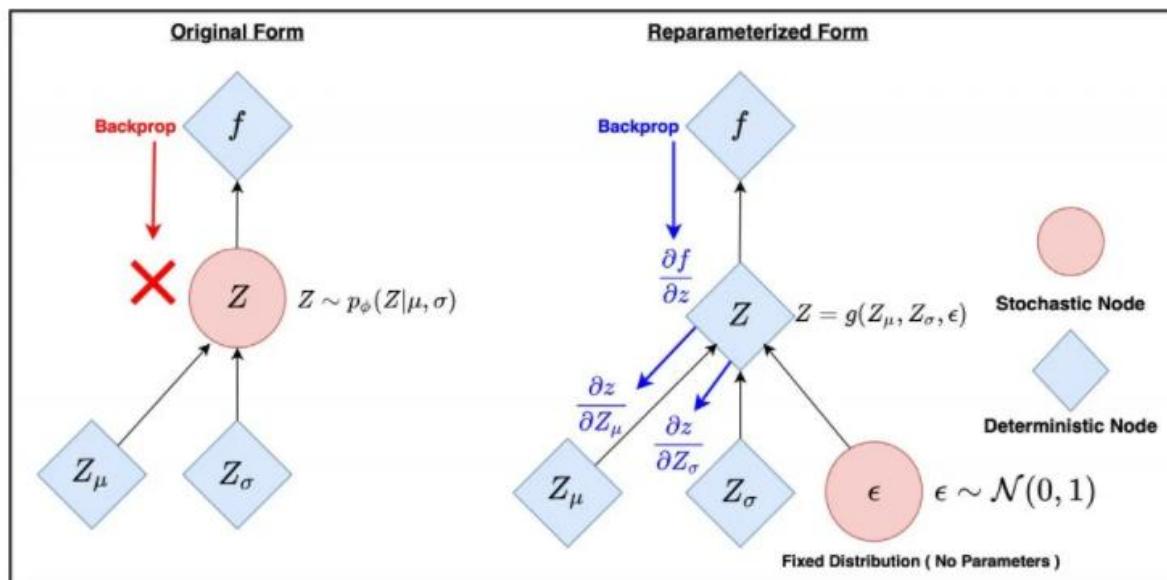
$$\mathcal{L}_{BCE} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Variational Autoencoder (VAE)

- **Re-parameterization Trick**

- Sampling prevents backpropagation and their training. (샘플링 $z \sim N(\mu, \sigma^2)$ 은 미분 불가능 연산)
- **Trick:** convert the random node Z to a deterministic node (Z_μ and Z_σ remain as the learnable parameters while still maintaining the stochasticity of the entire system via ϵ)

$$Z = Z_\mu + Z_\sigma^2 \odot \epsilon \quad (\text{Here, } \epsilon \sim \mathcal{N}(0, 1) \text{ and } \odot \text{ is element-wise multiplication.})$$



Variational Autoencoder (VAE)

- **KL Divergence**

- 두 확률 분포의 다른 정도 (= Relative Entropy)
- 두 분포가 Gaussian 일 경우 간략하게 표현 가능

$$-D_{KL}(q_\theta(z|x_i)||p(z)) =$$

$$\int \frac{1}{\sqrt{2\pi\sigma_q^2}} \exp\left(-\frac{(x-\mu_q)^2}{2\sigma_q^2}\right) \log \left(\frac{\frac{1}{\sqrt{2\pi\sigma_p^2}} \exp\left(-\frac{(x-\mu_p)^2}{2\sigma_p^2}\right)}{\frac{1}{\sqrt{2\pi\sigma_q^2}} \exp\left(-\frac{(x-\mu_q)^2}{2\sigma_q^2}\right)} \right) dz$$

$$= \int \frac{1}{\sqrt{2\pi\sigma_q^2}} \exp\left(-\frac{(x-\mu_q)^2}{2\sigma_q^2}\right) \times \\ \left\{ -\frac{1}{2} \log(2\pi) - \log(\sigma_p) - \frac{(x-\mu_p)^2}{2\sigma_p^2} + \frac{1}{2} \log(2\pi) + \log(\sigma_q) + \frac{(x-\mu_q)^2}{2\sigma_q^2} \right\} dz$$



$$q_\theta(z|x_i) \rightarrow \frac{1}{\sqrt{2\pi\sigma_q^2}} \exp\left(-\frac{(x-\mu_q)^2}{2\sigma_q^2}\right)$$

$$p(z) \rightarrow \frac{1}{\sqrt{2\pi\sigma_p^2}} \exp\left(-\frac{(x-\mu_p)^2}{2\sigma_p^2}\right)$$

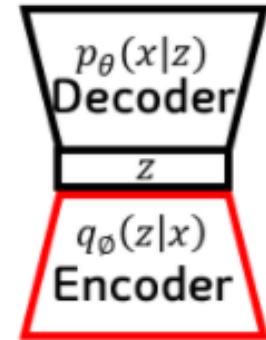
$$\sigma_p = 1 \text{ and } \mu_p = 0,$$

$$\begin{aligned} -D_{KL}(q_\theta(z|x_i)||p(z)) &= \log(\sigma_q) - \frac{\sigma_q^2 + \mu_q^2}{2} + \frac{1}{2} \\ &= \frac{1}{2} \log(\sigma_q^2) - \frac{\sigma_q^2 + \mu_q^2}{2} + \frac{1}{2} \\ &= \frac{1}{2} \left[1 + \log(\sigma_q^2) - \sigma_q^2 - \mu_q^2 \right] \end{aligned}$$

Variational Autoencoder (VAE)

- VAE Loss function

$$L_i(\phi, \theta, x_i) = \frac{-\mathbb{E}_{q_\phi(z|x_i)}[\log(p_\theta(x_i|z))]}{\text{Reconstruction Error}} + \frac{KL(q_\phi(z|x_i)||p(z))}{\text{Regularization}}$$



$$\mathcal{L} = -\sum_{j=1}^J \frac{1}{2} \left[1 + \log (\sigma_i^2) - \sigma_i^2 - \mu_i^2 \right] - \frac{1}{L} \sum_l E_{\sim q_\theta(z|x_i)} \left[\log p(x_i|z^{(i,l)}) \right]$$

$$(\theta^*, \phi^*) = \operatorname{argmin}_{(\theta, \phi)} \mathcal{L}(\theta, \phi)$$



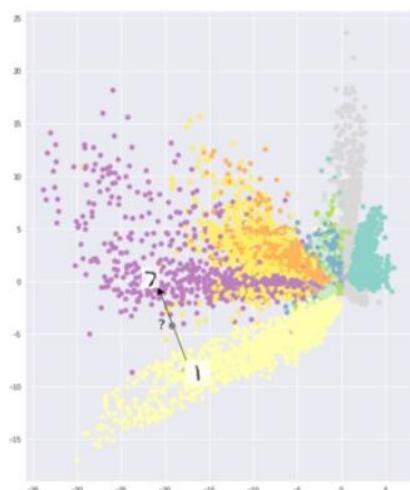
Cross entropy

L: number of samples
J: dim of latent vector z

Variational Autoencoder (VAE)

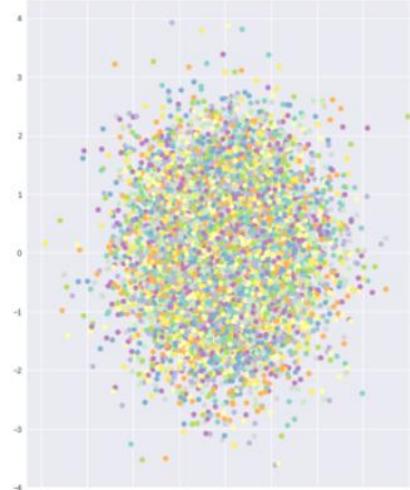
$$\mathcal{L}(x, \hat{x}) + \beta \sum_j KL(q_j(z|x) || N(0, 1))$$

Only reconstruction loss



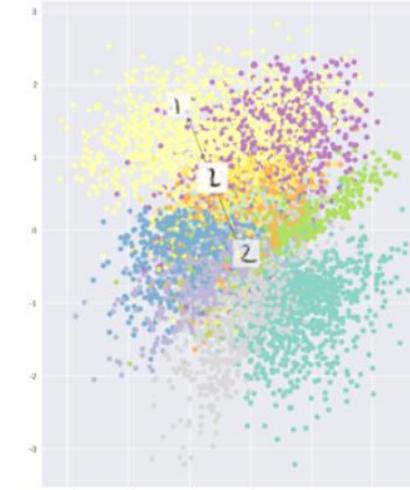
- Focus only on reconstruction loss, and allows the decoder be able to reproduce the original handwritten digits.
- there are areas in latent space which don't represent *any* of our observed data

Only KL divergence



- Focus only on KL divergence loss, and describe every observation using the same Gaussian (e.g. same characteristics)
- failed to describe the original data

Combination

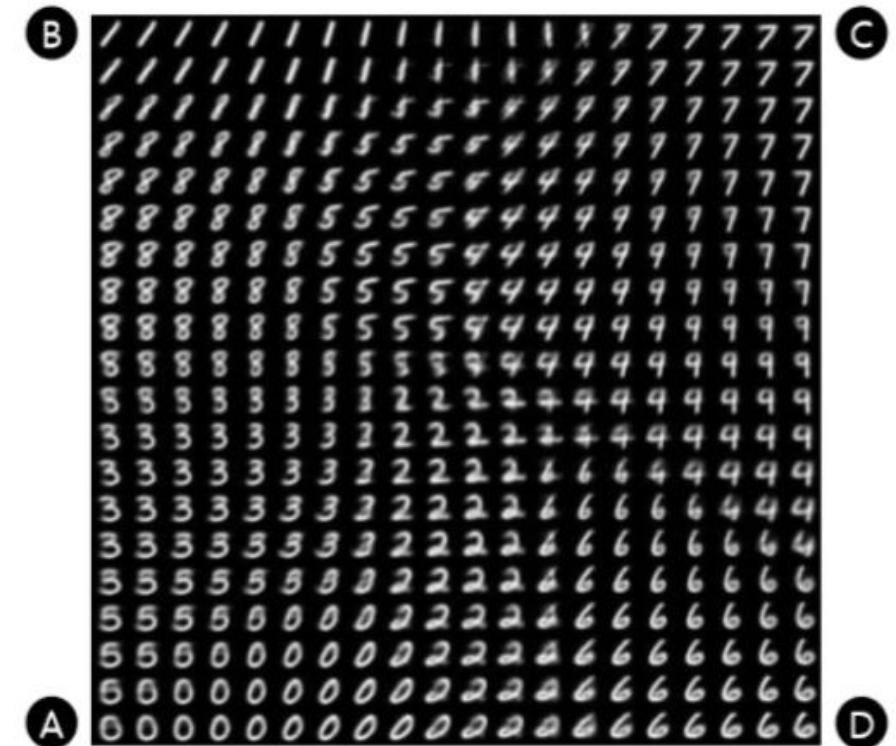
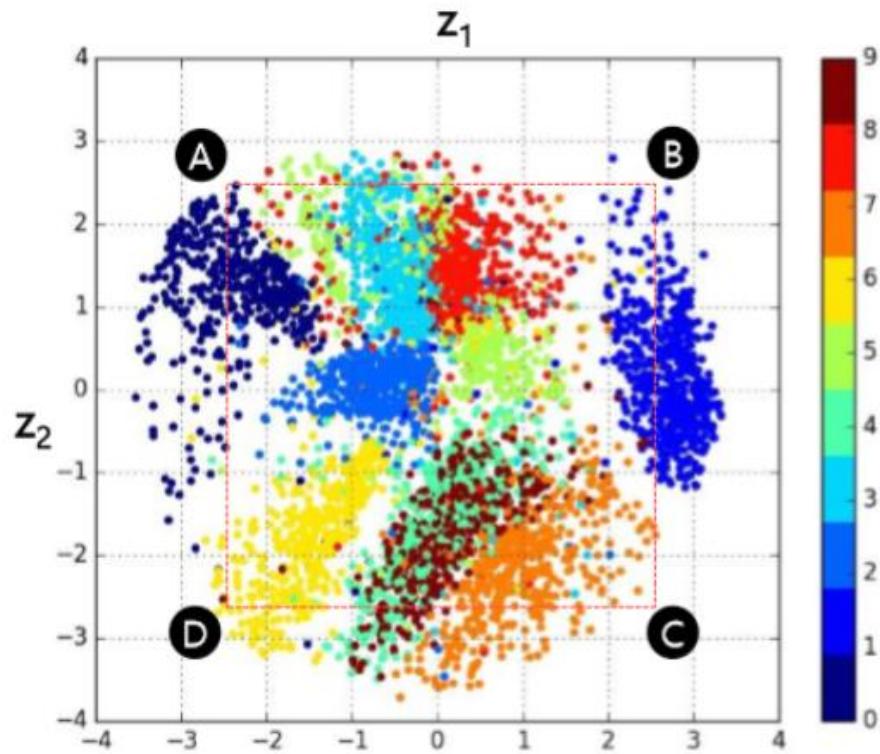


Variational Autoencoder (VAE)

- An example of well-trained VAE with MNIST dataset

$$z = \mu + \sigma \cdot \epsilon, \quad \epsilon \sim \mathcal{N}(0, 1)$$

Learned Manifold



학습이 잘 되었을 수록 2D공간에서 같은 숫자들을 생성하는 z들은 뭉쳐있고,
다른 숫자들은 생성하는 z들은 떨어져 있어야 한다.

Variational Autoencoder (VAE)

- **Another example**

- Gradually changing hair color and sunglasses to a face
- <https://github.com/comphree/variational-autoencoder>

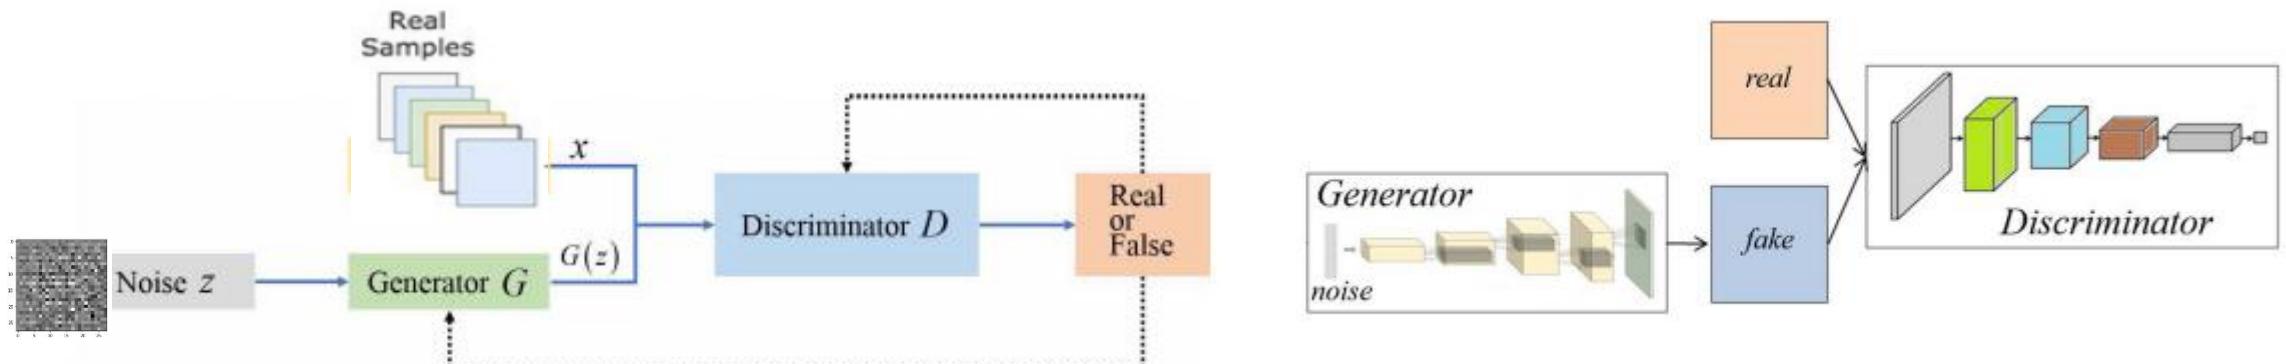
- ϵ 를 고정하면: 같은 입력 x 에 대해 항상 동일한 z 가 나오므로, 결과 이미지가 안정적으로 재현.
- ϵ 를 바꾸면: 같은 입력이라도 다른 잠재 벡터가 생성되어, 다양한 이미지 변형을 얻을 수 있다.
- ϵ 의 크기를 조절하면:
 - 작은 $\epsilon \rightarrow \mu$ 근처의 안정적인 샘플 \rightarrow 원본과 유사한 이미지
 - 큰 $\epsilon \rightarrow$ 잠재 공간의 더 먼 영역 \rightarrow 원본과 다른, 창의적인 이미지



GAN

- **Generative Adversarial Network (GAN)**

- Loss function: to make $\text{real}(D(x))$ close to 1, and $\text{fake}(D(G(z)))$ close to 0
- Discriminator: binary classifier (1: real, 0: fake)
- Two separated optimizers



$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim P_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim P_{\text{noise}}} \left[\log (1 - D(G(z))) \right]$$

\leftarrow

$$H_p(q) = -\frac{1}{N} \sum_{i=1}^N \underbrace{y_i \cdot \log(p(y_i))}_{y_i=1 \text{ 의 loss function}} + \underbrace{(1 - y_i) \cdot \log(1 - p(y_i))}_{y_i=0 \text{ 의 loss function}}$$

GAN – Loss function

1.1 판별자 (D)

- **역할:** 실제 데이터와 생성된 데이터를 구분.
- **목표:** 실제 데이터 x 에 대해 $D(x) = 1$, 생성 데이터 $G(z)$ 에 대해 $D(G(z)) = 0$.
- **손실 함수:**

$$\mathcal{L}_D = \frac{-\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]}{\text{실제 데이터 샘플 } x \text{ 에 대해 } D(x) \text{ 가 } 1 \text{에 가까워지도록 학습}} \quad \begin{matrix} \uparrow \\ \text{생성 데이터 } G(z) \text{ 에 대해 } D(G(z)) \text{ 가 } 0 \text{에 가까워지도록 학습} \end{matrix}$$

1.2 생성자 (G)

- **역할:** 실제처럼 보이는 데이터를 생성하여 D 를 속임.
- **목표:** 생성된 데이터 $G(z)$ 가 실제 데이터처럼 보이도록 $D(G(z)) = 1$ 로 만들기.
- **손실 함수:**

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

생성자는 판별자를 속이는 게 목표이므로 $D(G(z))$ 가 1에 가까워지도록 학습 (클래스 0 은 생성자 입장에서는 고려할 필요가 없음)

GAN – Loss function

$$\min_G \max_D \mathcal{L}(D, G)$$

$$\mathcal{L}(D, G) = \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Discriminator: $\max_D \mathcal{L}(D, G)$

Generator: $\min_G \mathcal{L}(D, G)$

실제 학습 (Minimize)

Discriminator Loss (\mathcal{L}_D):

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Generator Loss (\mathcal{L}_G):

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

생성자 (G) loss:

- 원래 minimax 공식:

$$L_G = \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

- Non-saturating Loss (더 안정적):

$$L_G = -\mathbb{E}_{z \sim p_z} [\log D(G(z))]$$

GAN

- **Generative Adversarial Network (GAN)**

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim P_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim P_{\text{noise}}} \left[\log (1 - D(G(z))) \right]$$

```
def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    total_loss = real_loss + fake_loss
    return total_loss
```

```
def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

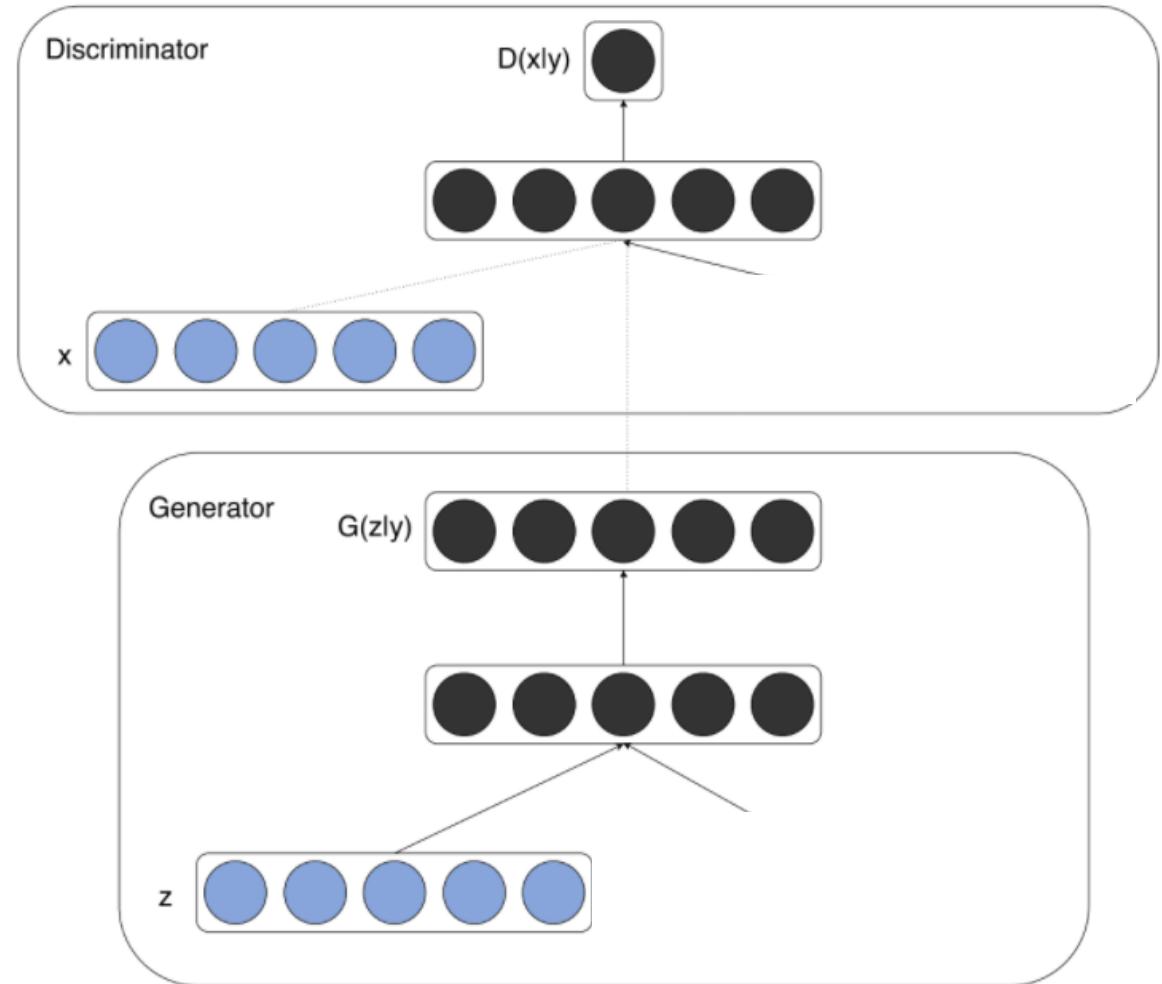
```
generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)
```

[ref: <https://www.tensorflow.org/tutorials/generative/dcgan?hl=ko>]

(Unconditional) GAN

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

↑
Expectation
↑
x is sampled from real data
Probability of D(real)
z is sampled from $N(0, 1)$
Probability of D(fake)
fake



Example of a Conditional Generator and a Conditional Discriminator in a Conditional Generative Adversarial Network.
Taken from Conditional Generative Adversarial Nets, 2014.

Conditional GAN (cGAN)

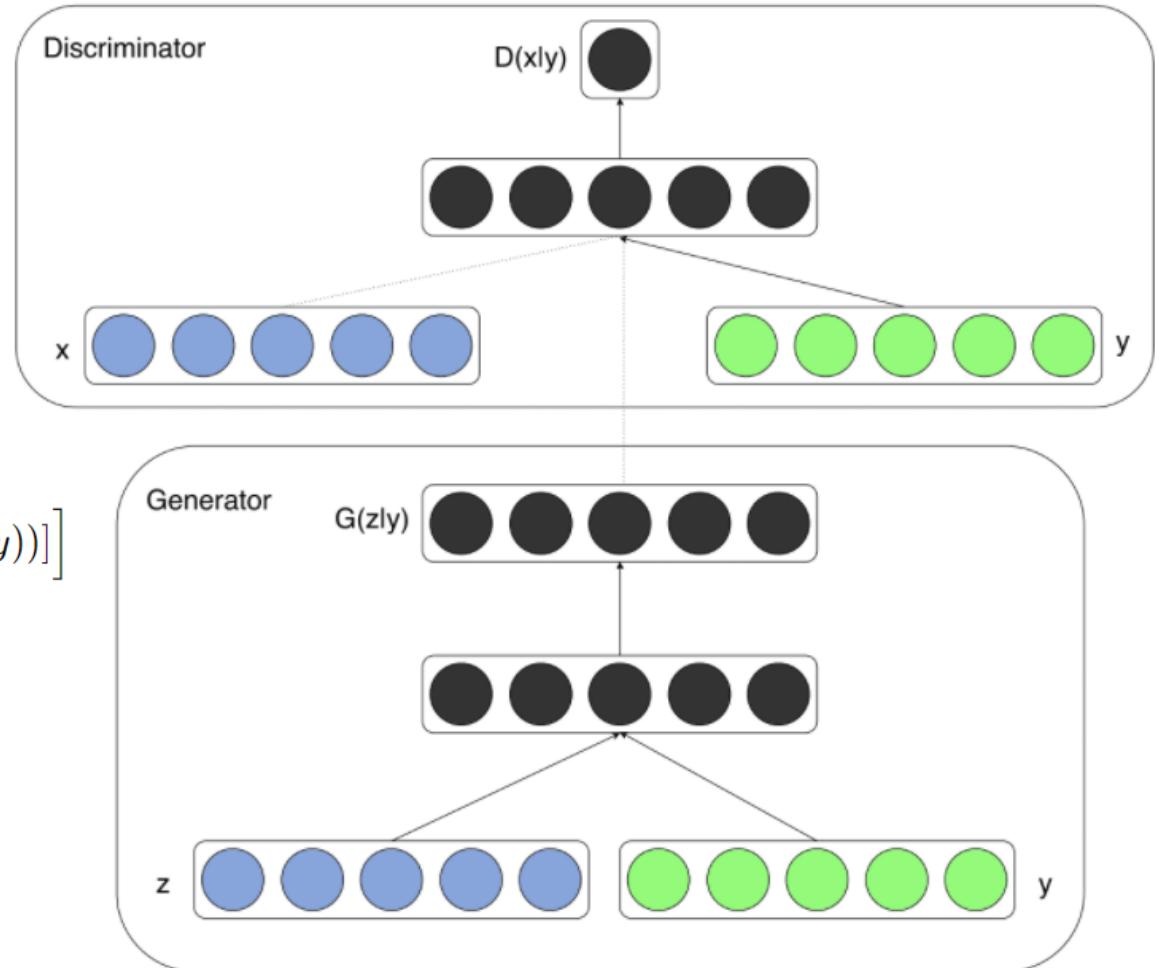
조건부 제어 생성

- y : 생성 과정에서 조건으로 제공되는 정보(예: 클래스 레이블).
- $G(z, y)$: 조건 y 에 따라 생성된 데이터 $G(z)$.
- $D(x, y)$: 데이터 x 와 조건 y 가 일치하는지 판별.

$$\mathcal{L}_{\text{cGAN}} = \min_G \max_D \left[\mathbb{E}_{x, y \sim p_{\text{data}}} [\log D(x, y)] + \mathbb{E}_{z \sim p_z, y \sim p_y} [\log(1 - D(G(z, y), y))] \right]$$

$D(x, y)$: 조건 y 에 맞는 진짜/가짜 판별

$G(z, y)$: "노이즈 z 와 조건 y 를 받아서, 조건에 맞는 이미지 생성"



Example of a Conditional Generator and a Conditional Discriminator in a Conditional Generative Adversarial Network.
Taken from Conditional Generative Adversarial Nets, 2014.